

Threads and Synchronization

CS 370 Spring 2019

Overview

- Review
- Deconstructing a Process
- Two Threads are Better Than One
- User level Threads vs Kernel Level Threads
- Convert Single Threaded to Multi Threaded
- POSIX Thread Library (pThread)

Review: Mutual Exclusion

- Correctness Property: Given two critical section executions for thread A and thread B: Mutual Exclusion property states that the i th execution of CS by A, CS^i_A , will **precede** the j th execution of CS by B, CS^j_B or CS^j_B will **precede** CS^i_A .
- Either $CS^i_A \rightarrow CS^j_B$ or $CS^j_B \rightarrow CS^i_A$ (they will not execute simultaneously)
- Deadlock satisfies this condition vacuously (no CS, no problems)
- Solutions to Mutual Exclusion also add:
- Deadlock free (always)
- Starvation free (sometimes)

Peterson's 2-Thread algorithm

Thread A

```
pre)  x: lockA = true;  
      y: turn = B;  
      z: while( lockB  
      w:      && turn != A );
```

```
post) x1: lockA = false;
```

Thread B

```
x: lockB = true;  
y: turn = A;  
z: while( lockA  
w:      && turn != B );
```

```
x1: lockB = false;
```

- What properties does this satisfy?
- Mutual Exclusion?
- Deadlock Free?
- Starvation Free?
- k-Bounded Waiting?

Bakery Algorithm

```
0 // declaration and initial values of global variables
1 Entering: array [1..NUM_THREADS] of bool = {false};
2 Number: array [1..NUM_THREADS] of integer = {0};
3
4 lock(integer i) {
5     Entering[i] = true;
6     Number[i] = 1 + max(Number[1], ..., Number[NUM_THREADS]);
7     Entering[i] = false;
8     for (integer j = 1; j <= NUM_THREADS; j++) {
9         // Wait until thread j receives its number:
10        while (Entering[j]) { /* nothing */ }
11        // Wait until all threads with smaller numbers or with the same
12        // number, but with higher priority, finish their work:
13        while ((Number[j] != 0) && ((Number[j], j) < (Number[i], i))) {
14        }
15    }
16 }
```

```
0 // declaration and initial values of global variables
1 Entering: array [1..NUM_THREADS] of bool = {false};
2 Number: array [1..NUM_THREADS] of integer = {0};
3
4 lock(integer i) {
5     Entering[i] = true;
6     Number[i] = 1 + max(Number[1], ..., Number[NUM_THREADS]);
7     Entering[i] = false;
8     for (integer j = 1; j <= NUM_THREADS; j++) {
9         // Wait until thread j receives its number:
10        while (Entering[j]) { /* nothing */ }
11        // Wait until all threads with smaller numbers or with the same
12        // number, but with higher priority, finish their work:
13        while ((Number[j] != 0) && ((Number[j], j) < (Number[i], i))) {
14        }
15    }
16 }
```

- What properties does this satisfy?
- Mutual Exclusion?
- Deadlock Free?
- Starvation Free?
- k-Bounded Waiting?

Processes

- A process is an instance of an executing program
- Include current values of pc, registers, variables.
- **Multiprogramming** is the ability to switch between multiple processes in a system (Concurrency)
- Each process gets its own entry in the Process Control Block (PCB).
- Each process gets its own address space (memory)
- Each process runs a thread of execution and the PCB saves the state of the CPU when performing a context switch.

An Alternative View of a Process


- Two Independent Concepts: **Resource grouping** and **Execution**
- Resources
 - Address space (code and data)
 - Open files
 - Child processes
 - Accounting information (stats for scheduling, etc.)
- Thread of Execution
 - PC
 - Registers
 - Stack
 - State (Ready, Running, Waiting, Terminated)
- Processes are used to group resources together, threads are entities scheduled for execution on the CPU.

Threads

- “Two threads are better than one, but ten threads without wit are good as none.” – John Heywood, 1546
- A “light weight process”
- Each thread shares the same address space and resources with the parent process.
- A process can have multiple threads of execution (threads) which can be scheduled independently of each other.
- Threads work together to complete the task of the process.
- Safety?
 - Processes are isolated from each other
 - Threads are created by a process, so they are assumed to be safe
 - No precautions are taken

Benefits of Using Threads

- An application may have multiple independent tasks to do at once
 - Threading simplifies this by decomposing each into separate task threads
- Faster to create and destroy threads than processes (10-100 times faster)
 - **Economy**
- Share an address space for fast inter-thread communication
 - **Resource Sharing**
- Better performance in the entire application
 - I/O bound threads don't block the entire application
 - Increased **Responsiveness**
- Multiprocessing (Parallelism)
 - **Scalability**



Single Thread to Multi Thread

- What things can go wrong?
- Global variables read and written by two or more threads simultaneously
- Can library calls support calls in parallel?
- Signals like alarms and interrupts
 - Who should receive keyboard interrupt? One thread, all threads?
- Stack management in user level threads


Moore's Law

- Gordon Moore in 2015:
- “I see Moore’s Law dying here in the next decade or so.”

Moore’s Law – The number of transistors on integrated circuit chips (1971-2016)

OurWorld
in Data

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are strongly linked to Moore's law.




Data source: Wikipedia (https://en.wikipedia.org/wiki/Transistor_count)

The data visualization is available at [OurWorldinData.org](https://www.ourworldindata.org). There you find more visualizations and research on this topic.

Licensed under CC-BY-SA by the author Max Roser.

Amdahl's Law

- Gives an upper bound for the *speedup* factor of adding N cores to a program
- $speedup \leq \frac{1}{S + \frac{1-S}{N}}$
- S is the portion of the program that must be run sequentially
 - $0 \leq S \leq 1$
- $1 - S$ is the parallel portion



User Level vs Kernel Level




Figure 2-16. (a) A user-level threads package. (b) A threads package managed by the kernel.

User Level Threads

- Each process needs its own thread table (contains pc, registers, state, stack)
- At least 10x faster than invoking the kernel
- User gets to choose proprietary scheduling algorithm
- Can decide when is a good time for a thread to block in terms of the program execution rather than using time slices
- **Problem:** Blocking System Calls
 - What happens when a thread calls *read()*?
- **Problem:** Page Faults Block
- **Problem:** Multiprocessing environments
 - Can user-level threads achieve “true parallelism” if multiple CPUs are available?

Kernel Level Threads

- Kernel keeps the thread table of all threads in the system
 - Keeps track of pc, registers, stack, state
- Thread creation and deletion invokes the kernel (slower)
- Subject to the kernel's scheduling algorithm
- May run another thread of the same process, or a different process's thread
- Resolves problems with system calls
- Can take advantage of multiprocessing!

Multithreading Models

- Many to One
 - Multiple user level threads supported by a single kernel level thread
- One to One
 - Each user level thread is supported by a kernel level thread
- Many to Many
 - Hybrid model, N user level thread supported by up to N kernel threads

Parallelizing a Problem

How to think about Threading

- Embarrassingly Parallel
 - Problem can be broken down without dependence on other threads
 - Rendering frames for display
 - Password Cracking
- Inherently Sequential Problems
 - Problem can not be broken down into threads
 - Iterative methods e.g. Newton's Method
- Matrix Multiplication, $C = A * B$
 - Fairly difficult to parallelize, can be done
 - Cache becomes important.
 - Use B transpose! Multiply rows in parallel
 - ~8x speedup with 4 threads on 4 core machine and matrix sizes $> 2000 \times 2000$
 - Superlinear speedup! Tomikj and Gusev paper, and Gusev and Ristov paper

Computers are Asynchronous

- Threads can experience sudden, unpredictable delays:
- Cache Misses
 - ~1-10 cycles
- Page faults (long)
 - ~1,000s of cycles
- Scheduling timeslice used up (really long)
 - ~1,000,000s of cycles
- These are common cases for why a thread is called “slow” or “lazy” in reasoning about parallel algorithms.

Race Conditions



Cntr++

Cntr++



[This Photo](#) by Unknown Author is licensed under [CC BY-ND](#)
[This Photo](#) by Unknown Author is licensed under [CC BY](#)

[This Photo](#) by Unknown Author
is licensed under [CC BY-SA](#)

Race Conditions

- When the behavior of a system is **dependent** on the sequence or timing of other uncontrollable events
- Ex: Two threads writing simultaneously to a shared counter
- Are race conditions always bad?
 - No!
 - Ex: Two threads read the same constant variable at the same time
- If the counter does not need to return unique values, then race condition is ok
 - Ex: random number generator's seed value in some contexts

Making a Counter

- `readVal = cntr.getAndIncrement()`
- Reads the value and increments it
 - Like `readVal = cntr++;`
- How can we avoid a race condition on `cntr`?
- `cntr.getAndIncrement(){`
 - `temp = cntr;`
 - `cntr = temp + 1;`
 - `return temp;`
- `}`

Solutions



[This Photo](#) by Unknown
Author is licensed under
[CC BY-SA](#)



[This Photo](#) by Unknown Author is
licensed under [CC BY-SA](#)

Making a Counter Work

- Intuition: How should a counter work?
- Correctness Conditions (Safety Properties)
 - Nothing bad will ever happen
- Progress Conditions (Liveness Properties)
 - Something good happens eventually
- What does a **correct** counter do?
- What kind of **progress** can a thread make?

Making a Counter Work

- Specification:
 - Must return unique values to all calling threads
 - Must be monotonically increasing
- No deadlock

Counter Solutions

- Do nothing
 - Hope that nothing bad happens
 - Allows race condition, does not solve the problem
- Lock access to function getAndIncrement
 - **Pro:** Solves race condition
 - Uses mutual exclusion to guarantee correctness
 - Algorithm for mutual exclusion will have some progress guarantee
 - Deadlock free, starvation free, etc.
 - **Con:** Added overhead of mutual exclusion algorithm
- Glue instructions together into one indivisible unit
 - **Pro:** Solves race condition
 - Use stronger primitives to avoid using locks
 - Hardware R-M-W operations
 - **Con:** Using stronger instructions too much can slow down performance



[This Photo](#) by Unknown
Author is licensed under [CC BY-SA](#)

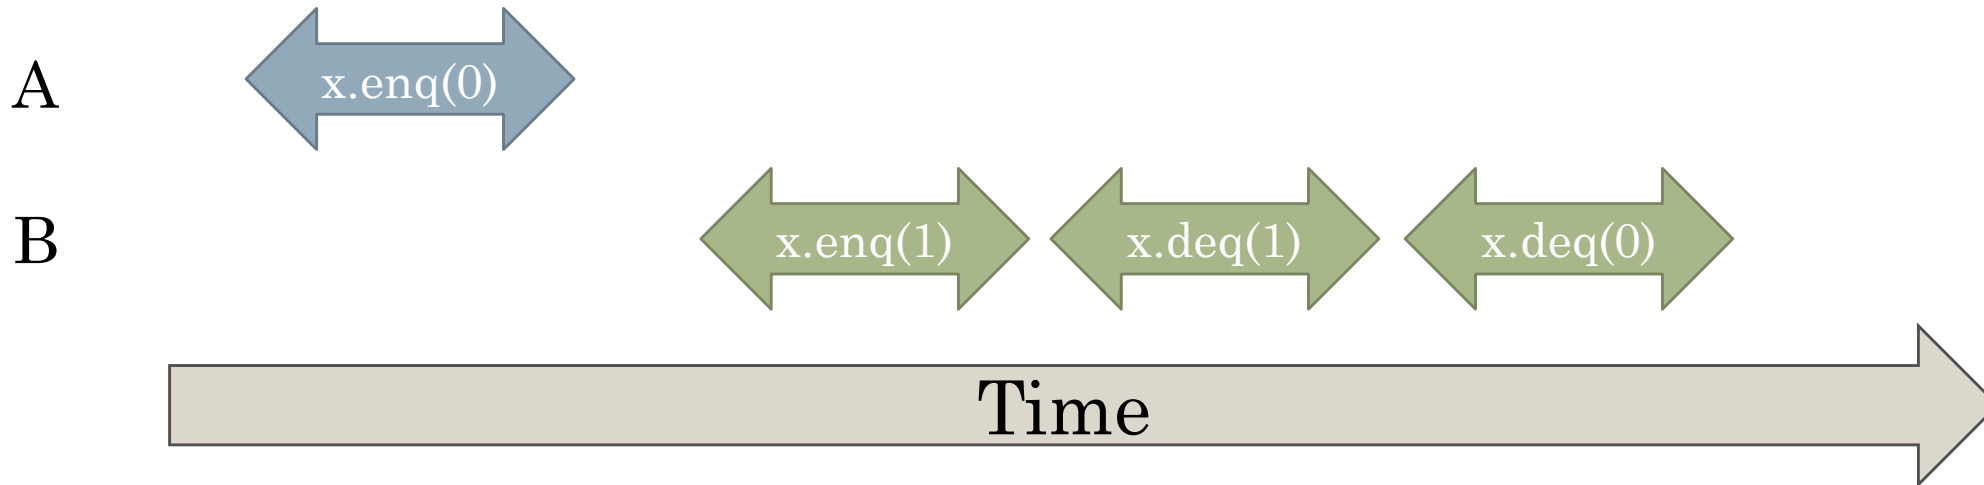


[This Photo](#) by Unknown Author
is licensed under [CC BY-SA](#)



Correctness Condition

- **Program Order:** The order of instructions in which a program is written
- **Sequentially Consistent:**
 - Each thread must take effect in program order
 - Can interleave calls of threads in any order we want.
- C11 Atomic class guarantees this!



Progress Conditions

- These guarantee some progress will be made
 - “Liveness”
- Deadlock Free
 - Some thread trying to acquire the lock will succeed
- Starvation Free
 - All threads trying to acquire the lock will eventually succeed
- Lock Free
 - Some thread calling a method will return
- Wait Free
 - All threads that call a method will return


Locks

- Mutexes
 - `pthread_mutex_t`
 - Lock
 - One thread can acquire the mutex at a time
 - Unlock
 - Owner must unlock the lock
- Semaphores
 - `pthread_semaphore_t`
 - Initialize the value to an integer
 - Wait
 - Wait for available token (value > 0)
 - Signal/Post
 - Add a token
- Is a semaphore initialized to 1 the same as a mutex?



Atomics R-W

- Atomic means unable to be divided
- Can read or write to a location without being interrupted
- `#include <atomic.h>`
- `atomic_load(atomic_type* A);`
- `atomic_store(atomic_type* A, type V);`
- Sequentially Consistent Correctness Condition
 - Each thread must take effect in program order
 - Can interleave calls of threads in any order we want.
- Are atomic variables lock free?



Stronger Primitives R-M-W

- The Lock-free and Wait-free Toolbox
- Test-and-Set
- Fetch-and-Add
- Compare-And-Set
- RISC: Load-Link (LL) and Store-Conditional (SC)



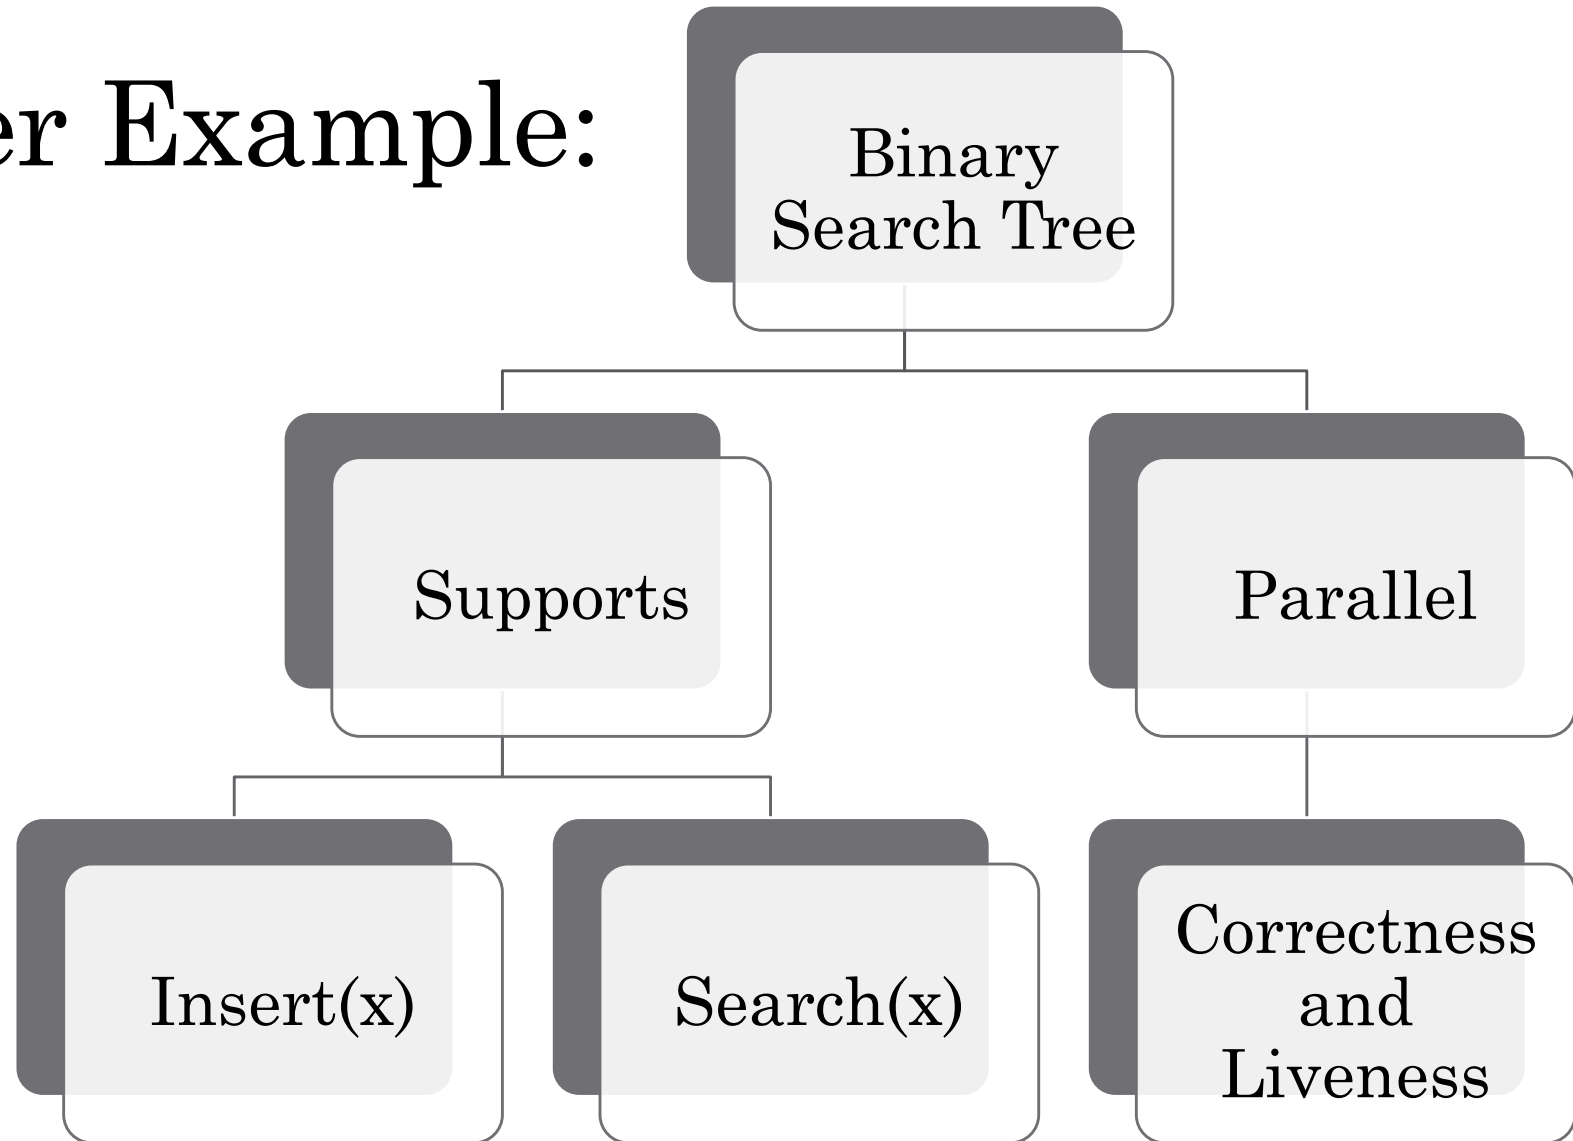
Examples

Another Example: Primes

- You want to find all the primes between 1 and N :
- Sequentially
- In Parallel with P processors
- Correctness?
- Liveness?



Another Example:



BST Search(x, root)

- Search(x, root){
 1. itr = root
 2. while(itr != nil){
 3. if(itr.item == x)
 4. return itr
 5. if(itr.item < x)
 6. if(itr.leftChild == nil)
 7. return itr
 8. itr = itr.leftChild
 9. else
 10. if(itr.rightChild == nil)
 11. return itr
 12. itr = itr.rightChild
 13. }
 14. return nil}

BST Insert(x, root)

- Insert(x, root){
 1. while(true){
 2. node = Search(x, root)
 3. if node.item == x
 4. return false
 5. if(CAS(node, nil, node(k) == true)
 6. return
 7. root = node}

5 Challenges

- Identifying Tasks
 - What ideally independent tasks can be run in parallel?
- Balance
 - Does each thread have nearly equal work?
- Data Splitting
 - How should data accessed and changed by threads be split?
- Data Dependency
 - If two or more threads access/change the same data or depend on data from another, synchronization may be necessary.
- Testing and Debugging
 - Due to asynchrony, threads can run at different times and the interleaving of procedures can occur in many different ways

Types of Parallelism

- Data Parallelism
 - Share data among threads, each thread does the same operation
 - Ex: Two threads sum half the data each and combine
- Task Parallelism
 - Each thread does different tasks
 - Ex: Word processor with different threads for keyboard input, spell checking, saving backups periodically
 - Ex: One thread computes the mean, another thread computes the median
 - Ex: Web server with dispatch and worker threads

Lock or Not?

- Sometimes a clever way to think about the problem can yield a good wait-free or lock-free solution
- Sometimes good wait-free and lock-free solutions are simple!
- Don't be quick to throw locks at a problem!
- Other times, we can't eliminate all dependencies without incurring large overhead
 - Use locks!
 - Try to keep Critical Sections small.

POSIX Threads

POSIX Thread Library (pThread)

- A specification **not** an implementation.
- `#include <pthread.h>`
- Create a thread with:
- **`int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void *(*start_routine) (void *), void *arg);`**
- Exit the thread with:
- **`void pthread_exit(void *retval);`**
- Offers synchronization objects too:
 - `pthread_mutex_t`
 - `pthread_semaphore_t`
 - `pthread_cond_t`
- See handout for Project 5 with detailed syntax

pThread Continued

- Main function typically will create the threads
- “With great power comes great responsibility”
- Function that creates the threads should ensure they finish before terminating itself
- Main will **wait** for the threads to finish
- Then main will **join** with the threads
- **int pthread_join(pthread_t *thread*, void ***retval*);**
 - If thread has called exit, then this call just gets return val and completes
 - Else, caller will **wait** until thread has called exit and then get return val

pthread_mutex_t

- `pthread_mutex_init(pthread_mutex_t *mutex, const pthread_mutexattr_t *attr);`
 - Initialize the mutex
- `pthread_mutex_destroy(pthread_mutex_t *mutex);`
 - Destroy the mutex
- `int pthread_mutex_lock(pthread_mutex_t *mutex);`
 - Lock the mutex, only owner can unlock
- `int pthread_mutex_unlock(pthread_mutex_t *mutex);`
 - Unlock mutex

pthread_semaphore_t

- `int sem_init(sem_t * sem, int pshared, unsigned int value);`
 - Initialize semaphore *sem* to *value*
 - *pshared* flag is **1** for between process semaphore and **0** for thread-access only
- `int sem_destroy(sem_t *sem);`
 - Delete the semaphore and mark memory for reallocation
- `int sem_post(sem_t *sem);`
 - Signal, add one to the value of semaphore *value*
- `int sem_wait(sem_t *sem);`
 - if *value* > 0, then *value*--
 - if *value* == 0, wait

pthread_cond_t

- `int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex);`
 - Wait for a signal on the *cond* using *mutex*
- `int pthread_cond_signal(pthread_cond_t *cond);`
 - Wakes up one thread arbitrarily that is waiting on *cond*
- `int pthread_cond_broadcast(pthread_cond_t *cond);`
 - Wakes up **all** threads that are waiting on *cond*

References

- Tannenbaum, Bos “Modern Operating Systems” Fourth Edition
 - Chapter 2
- Abraham, Silberschatz “Operating System Concepts” 9th Edition
 - Chapter 4
- https://en.wikipedia.org/wiki/Amdahl's_law#/media/File:AmdahlsLaw.svg
- Tomikj, Nikola, and Marjan Gusev. "Parallel Matrix Multiplication." *2018 41st International Convention on Information and Communication Technology, Electronics and Microelectronics (MIPRO)* (2018): 0204-209. Web. <https://doi.org/10.23919/MIPRO.2018.8400039>
- Gusev, Marjan, and Sasko Ristov. "A Superlinear Speedup Region for Matrix Multiplication." *Concurrency and Computation: Practice and Experience* 26.11 (2014): 1847-868. Web. <https://doi-org.ezproxy.library.unlv.edu/10.1002/cpe.3102>
- <https://spectrum.ieee.org/computing/hardware/gordon-moore-the-man-whose-name-means-progress>