

# Operating System

## CS 370

### Spring 2019

#### Project 7

Due Date: 11:59 PM, April 23<sup>rd</sup> 2019  
*Must be submitted before deadline*

Total Points: 100  
Graduate Assistant: Pradip Maharjan  
Office Hours: MoWe 12 PM – 2PM  
Teaching Assistant: Luis Maya Aranda  
Office Hours: TuTh 10 AM – 12 PM

**Objective:** Become familiar with the memory system in xv6 and implement lazy page file allocation.

**Introduction:**

In xv6, new memory pages are allocated to a process as soon as they are requested. This is inefficient because it creates more overhead for the operating system, even if the pages are not used after they are requested. Therefore, many systems use a lazy scheme to only allocate a page once the first access occurs. This delays a significant amount of the page allocation work until the page is actually needed. In this exercise, you will implement lazy page file allocation in xv6.

**Resources:**

The following resources provide more in-depth information regarding **xv6**. They include the **xv6** reference book, source code (pdf), and a tutorial on running and debugging.

1. **xv6** Reference Book: <https://pdos.csail.mit.edu/6.828/2016/xv6/book-rev9.pdf>
- xv6** Source Code PDF: <https://pdos.csail.mit.edu/6.828/2016/xv6/xv6-rev9.pdf>
- Running and Debugging Tutorial: <http://zoo.cs.yale.edu/classes/cs422/2010/lec/l2-hw>

**Project:**

Complete the following steps.

- Become familiar with memory management in xv6
- Modify xv6 to implement lazy page file allocation

**Submission**

When complete, submit:

- Zip the entire **xv6** folder and submit (not qemu).
- A copy of the **zipped xv6 folder** via the class web page (assignment submission link) by class time (1:00 PM).

**Submissions received after the due date/time will not be accepted.**

### **Background**

Before implementing lazy page file allocation, it is important to familiarize yourself with the existing memory system. First, review the process entry in `proc.h`. Note that the page directory (the per-process page table), `pgdir`, and the size of the process's memory, `sz`, are stored in each process's entry in the process control block. In the existing system, a user program can request more memory for its program by calling `malloc()` which calls `sbrk()` (implemented in `sysproc.c` as `sys_sbrk()`) which then calls `growproc()`. See the implementation of `sbrk()` in `sysproc.c` and the implementation of `growproc()` in `proc.c`. The function `growproc()` in `proc.c` calls `allocuvm()` to allocate `n` bytes to the calling process's memory. Familiarize yourself with the implementation of `allocuvm()` in `vm.c`. The function `allocuvm()` allocates all the pages up until the new memory size is reached. The `exec()` function also calls `allocuvm()` to load a new process, see `exec.c`.

### **Modify sbrk()**

When `sbrk()` is called, it calls `growproc()`, which in turn calls `allocuvm()` which will allocate all the pages requested immediately. To switch to lazy page allocation, our first task is to remove the call to `growproc()` from `sys_sbrk()` in `sysproc.c`. Instead, have `sbrk` just increment the size of the process's size by `n` bytes, and save the allocation for later.

If you run `xv6` with this change alone, you will get a page fault error when you run a user program. Try it. You can use `traps.h` to check what the trap error corresponds to.

### **Catch Page Faults and Allocate**

Recall from an earlier assignment that trap errors are generated in `trap.c`. You can test the type of trap that has occurred and see if it is a page fault by checking the **`trapno`**. Use the code that allocates a page present in `allocuvm()` to allocate the page that the page fault occurred on. You can use the function **`rcr2()`** (read control register 2, see CR2 in x86 architectures for more information) to get the address that caused the page fault. Note that you will have to call `mappages()`, which is not defined in `proc.c`. Add the function prototype to the beginning of `trap.c` to allow the program to find the function:

```
int mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm);
```

### **Modify copyuvm() to make fork() functional**

The `fork()` system call uses `copyuvm()` in `vm.c` to copy the parent's heap memory. However, a parent may not have allocated all of the pages that it has requested at the time of `fork()`. This will cause a page not present error when forking. Modify `copyuvm()` in `vm.c` so that it skips copying pages that are not present. *Hint: `PTE_P` bitwise OR'ed with the page table entry address to determine if a page is present.*

### **Modify mappages() to Skip Present Pages**

Currently, the `mappages()` function will cause a panic when a page is already present. Remove the panic in `mappages()` to allow your benchmark program to function.

### **Testing Your Program**

Once you have all of the above changes made, test your program using all of the programs you have created this semester. `Xv6` should be functional with all of these programs.